



# Machine Learning

## Chapter 5: Nearest Neighbor

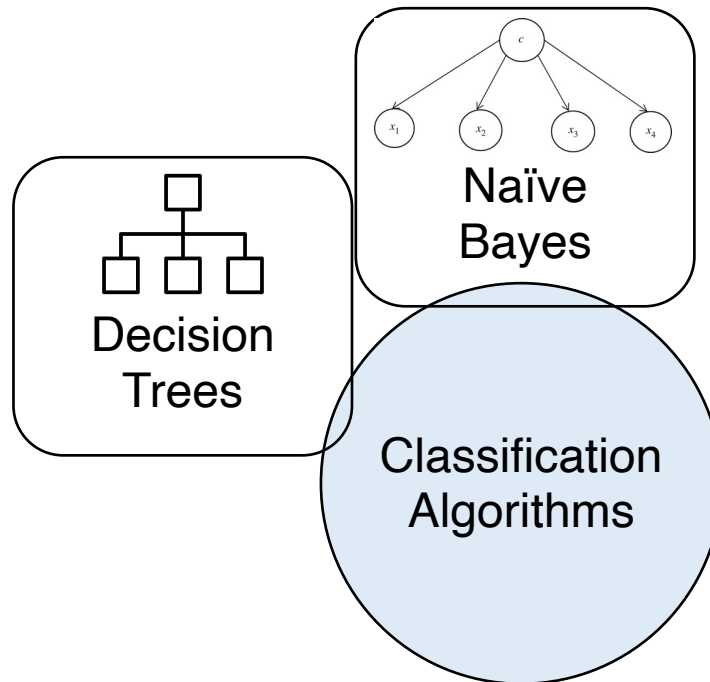
Fall 2022

Instructor: Xiaodong Gu

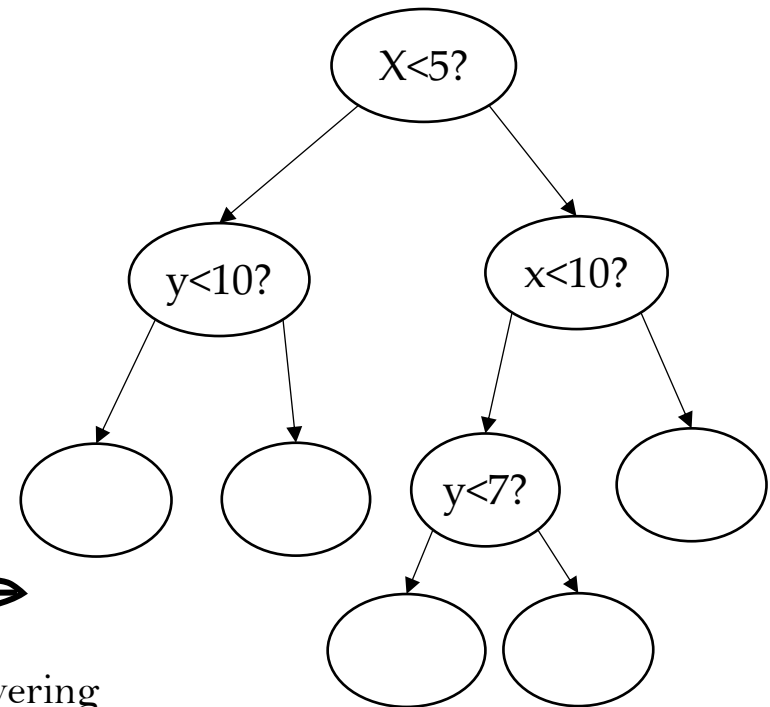





# The family of classification



# Review: Decision Trees




Combines a set of different linear models, each covering a region of the input space disjoint from others.

# Review: Bayes



Modeling the generation of data using probabilities?



have clear patterns of **probabilistic distributions** and **dependences**


# Do we really need a model ?



Data



Model



What if we directly use the data ?

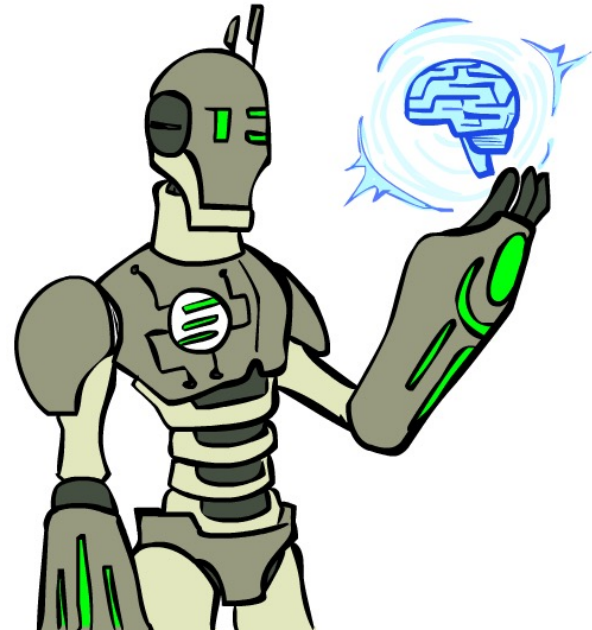
# Today

---



Let's learn a super easy (naïve) method for classification.

- K-Nearest Neighbor




# A Simple Analogy...

---



- Tell me about your friends (who your neighbors are).  
Then I will **tell you who you are**.



# K-Nearest Neighbors

Instance-based learning, also called lazy learning.

- simply stores the training instances without learning a model.
- whenever we have new data to classify, we find its K-nearest neighbors from the training data.




**K-Nearest Neighbor (KNN)** is a simple algorithm that stores all the available cases and classifies the new data or case based on a similarity measure.



# K-NN Classification



- Classified by “MAJORITY VOTES” from neighbor classes.
- An object is classified to the most common class amongst its  $k$  nearest neighbors **measured by a “distance” function**



How to determine whether an object falls into the  $k$ -nearest neighbors?



# Distance Measures

---

- **Euclidean Distance** 欧几里得距离

$$D(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

where  $\mathbf{x} = (x_1, x_2, \dots, x_n)$  and  $\mathbf{y} = (y_1, y_2, \dots, y_n)$  represent the  $n$  attribute values of two records.

doesn't work well in high dimensions and for categorical variables because it ignores the similarity between attributes.

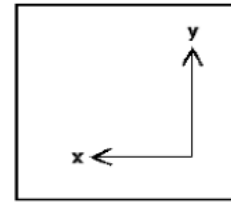


# Distance Measures

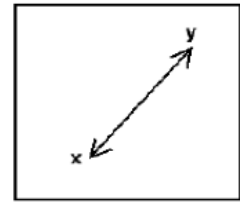
---

- Manhattan Distance (a.k.a. city block distance)

$$D(x, y) = \sum_{i=1}^n |x_i - y_i|$$



**Manhattan**



**Euclidean**

$$|x_1 - x_2| + |y_1 - y_2|$$


- Minkowski Distance

$$D(x, y) = \left( \sum_{u=1}^n |x_u - y_u|^p \right)^{\frac{1}{p}}$$

# Distances



## Example



point	x	y
p1	0	2
p2	2	0
p3	3	1
p4	5	1

Euclidean Distance Matrix

	p1	p2	p3	p4
p1	0	2.828	3.162	5.099
p2	2.828	0	1.414	3.162
p3	3.162	1.414	0	2
p4	5.099	3.162	2	0

# Normalization

---



- Standardize the range of independent variables (features of data)

**Z-score normalization:** rescale the data so that the mean is zero and the standard deviation from the mean (standard scores) is one.

$$X_{norm} = \frac{X - \mu}{\sigma}$$


**min-max normalization:** scale the data to a fixed range between 0 and 1.

$$X_{morm} = \frac{X - X_{min}}{X_{max} - X_{min}}$$



# Similarity vs. Distance

- **Similarity**: numerical measure of how alike two data objects are.
  - Is higher when objects are more alike.
  - Often falls in the range  $[0, 1]$ .
- **Cosine Similarity**



$$\cos(\theta) = \frac{a \bullet b}{||a|| \times ||b||}$$

$$= \frac{(x_1, y_1) \bullet (x_2, y_2)}{\sqrt{x_1^2 + y_1^2} \times \sqrt{x_2^2 + y_2^2}}$$

$$= \frac{x_1 x_2 + y_1 y_2}{\sqrt{x_1^2 + y_1^2} \times \sqrt{x_2^2 + y_2^2}}$$

$$\cos(d_1, d_2) = \begin{cases} 1: \text{exactly the same} \\ 0: \text{orthogonal} \\ -1: \text{exactly opposite} \end{cases}$$

# Quiz

---



$$\cos(d_1, d_2) = (d_1 \cdot d_2) / \|d_1\| \|d_2\|$$

$$d_1 = [3 \ 2 \ 0 \ 5 \ 0 \ 0 \ 0 \ 2 \ 0 \ 0]$$

$$d_2 = [1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 2]$$

$$d_1 \cdot d_2 = 3*1 + 2*0 + 0*0 + 5*0 + 0*0 + 0*0 + 0*0 + 2*1 + 0*0 + 0*2 = 5$$

$$\|d_1\| = (3*3 + 2*2 + 0*0 + 5*5 + 0*0 + 0*0 + 0*0 + 2*2 + 0*0 + 0*0)^{0.5} = (42)^{0.5} = 6.481$$

$$\|d_2\| = (1*1 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 1*1 + 0*0 + 2*2)^{0.5} = (6)^{0.5} = 2.245$$

$$\cos(d_1, d_2) = .3150$$

# The Algorithm

---




## Algorithm

---

1. Determine parameter  $K$
2. Choose a sample from the test data that needs to be classified and compute its distance to all the training examples.
3. Sort the distances obtained and take the  $k$ -nearest data samples.
4. Assign the test class to the class based on the majority vote of its  $k$  neighbors. (根据选出来的 $k$ 个neighbor中占多数的那个class作为当前这个data sample的预测类型结果)



# The Algorithm





---




What is the best value of  $K$  to use?

# Decision Boundary

## Voronoi Tessellation (沃罗诺伊分割)

- Partition the space into areas that are nearest to any given point
- Boundary: points at the same distance from two different training examples.

**Decision Boundary:** boundaries that separates two different classes.

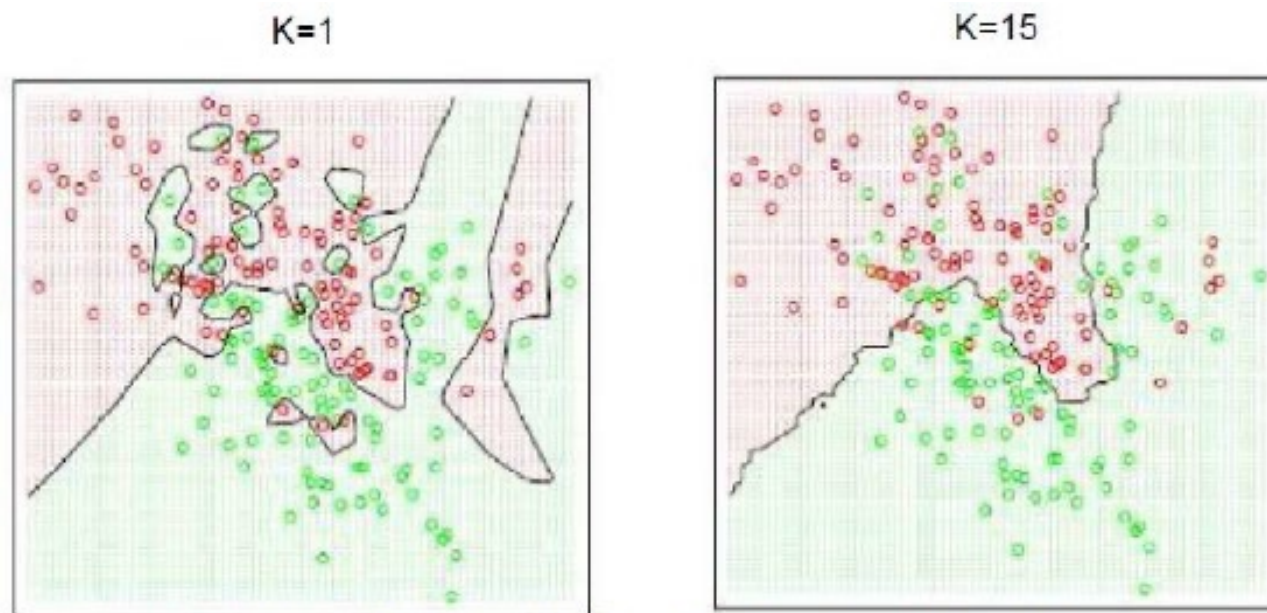


With large number of examples and possible noise in the labels, the decision boundary can become nasty!

# Effect of K

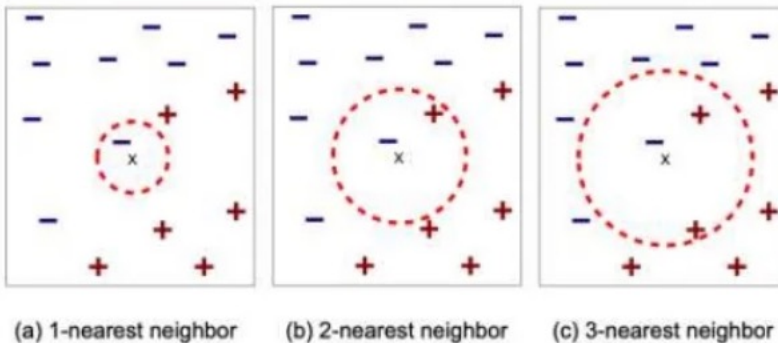
---

- Larger  $K$  produces smoother boundary effect
- When  $K=N$ , always predict the majority class



Figures from Hastie, Tibshirani and Friedman (Elements of Statistical Learning)

**Discussion:** which model is better between  $K=1$  and  $K=15$ ? Why?




### How to choose K?

- If K is too small, efficiency is increased but becomes susceptible to noise.
- Larger K works well. But too large K may include majority points from other classes, but risk of over-smoothing classification results

# Try it yourself

---



<http://vision.stanford.edu/teaching/cs231n-demos/knn/>

# Pros and Cons

---



## Advantages:

- Simple to understand, explain, and implement,
- No effort for training,
- New data can be added seamlessly without hampering the model accuracy

# Pros and Cons

---



## Disadvantages:

- Does not scale with large data sets (calculating distance is computationally expensive)
- Highly susceptible to the curse of dimensionality
- Large storage requirements
- Data normalization is required





---



# Tutorial: KNN with Python

# What's Next?

## Logistic Regression

Classification by a discriminative function.

- Linear functions
- Differentiable!



WHAT'S  
NEXT?