



Machine Learning

Lecture 9: Application - Word Embedding

单词嵌入

Fall 2022

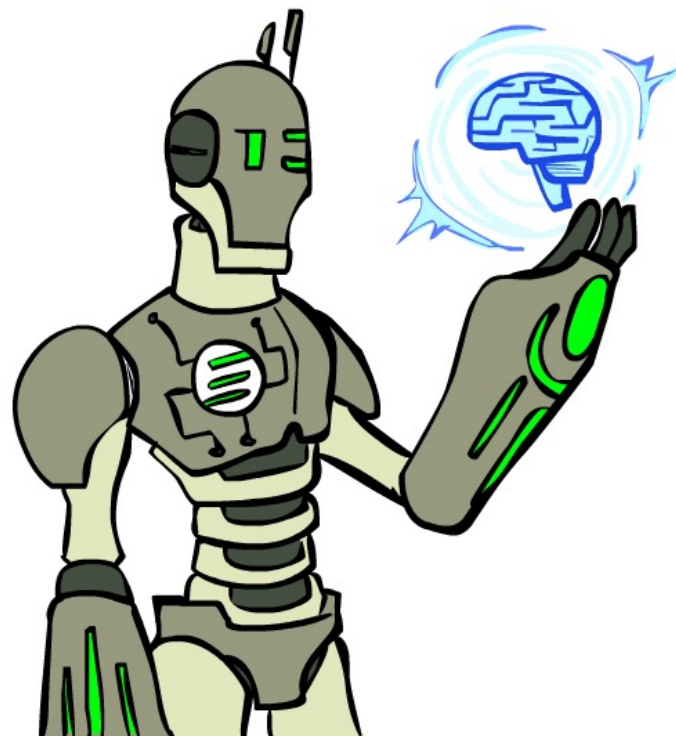
Instructor: Xiaodong Gu



Today



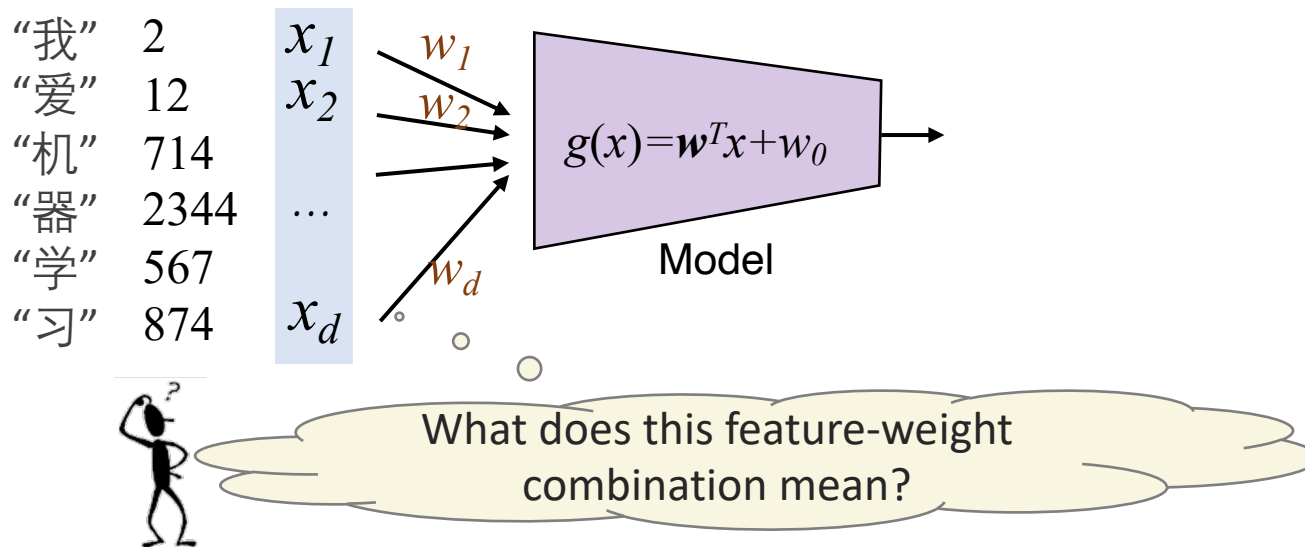
- Word Embedding (word2vec)
an application of **neural networks**
for distributed representations of words
- Language Model 语言模型
a key concept in NLP that is
important in the follow-up lectures.



Before learning word embedding...



- How to represent words in computers?
 - discrete-symbols (ASCII, id, etc.)
 - e.g., “hotel”, “conference”, “motel”, 1203, ...
 - a **localist** representation
- How to represent words in machine learning models?



Words as Vectors



- **Words** can be represented by **one-hot** vectors:

One-Hot Encoding

Means one 1, the rest 0s

hotel = [1 0 0 0 0 0 0 0 0]

flower = [0 1 0 0 0 0 0 0 0]

tree = [0 0 0 1 0 0 0 0 0]

motel = [0 0 0 0 0 0 0 1 0]

elephant = [0 0 0 0 0 0 0 0 1]

Vector dimensionality =
number of words in
vocabulary (e.g., **500,000**)

直接使用one-hot vector表示word的缺点：
1. 过多的0，向量过于稀疏
2. 只含有1与0，不好获取单词之间的关联性
以及相似性(因为两两正交)



How can we **capture relationships**
(similarity) between words?

Problems of One-hot Encoding



Example: in web search, if user searches for “Minhang motel”, we would like to match documents containing “Minhang hotel”.

But

motel = [0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]
hotel = [0 0 0 0 0 0 0 1 0 0 0 0 0 0 0]

orthogonal.

There is no natural notion of **similarity** for one-hot vectors!

Solution:

- Use **WordNet**'s list of synonyms?
- But it is well-known to fail badly: incompleteness, etc.



Instead: learn to encode similarity in the vectors themselves

Distributed Word Representation



- Represent words as **low-dimensional dense vectors** that can reflect their semantic similarities. dense : 密集的, 稠密的

One-Hot Encoding

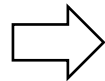
hotel = [1 0 0 0 0]

flower = [0 1 0 0

0]
business = [0 0 1 0 0]

motel = [0 0 0 1 0]

elephant = [0 0 0 0 1]



Word Embedding

● hotel

● motel

dog

elephant

cat

● tree

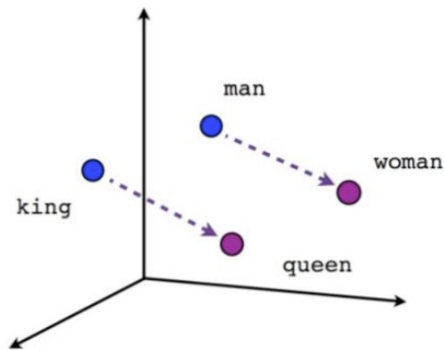
● flower

Note: word vectors are sometimes called **word embeddings** or word representations. They are **distributed** representations.

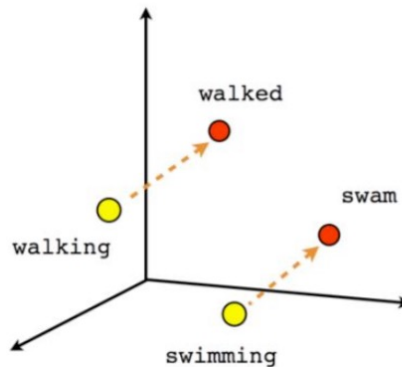
Why Word Embeddings?



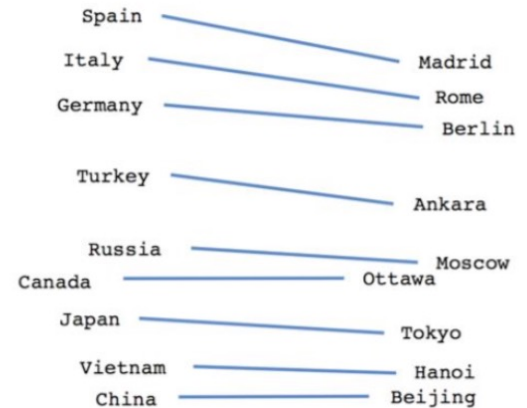
- Can capture the rich relational structure of the lexicon.



Male-Female



Verb tense



Country-Capital



How to obtain word embeddings?

How to obtain word embeddings?



- A word can be understood by **its context**

“A bottle of *tesgüino* is on the table”

“Everybody likes *tesgüino*”

“*Tesgüino* makes you drunk”

“We make *tesgüino* out of corn”



“You shall know
a word by the
company it
keeps” (J. R. Firth
1957)

- From context words we can guess **tesgüino** means an alcoholic beverage such as beer.

- Intuition for an algorithm:



Two words are similar if they have similar word contexts

idea1：根据单词的上下文来判断单词可能的语义，并根据上下文来判断单词的相似性。

How to Exploit the Context?



Counting-based Approach

two words have similar vectors if they frequently **co-occur** in the same context (sentence, document, etc.).

e.g. Glove:

Prediction-based Approach

train **neural networks** to predict a word (vector) **given its context words (vector)**.

e.g., word2vec



Counting based: the **vector space model**

(向量空间模型)

Vector Space Model



- The cornerstone technology in **information retrieval**.
- **Term-Document Matrix**

Each cell is the count of word t in document d

	d_1	d_2	d_3	d_4	d_5
ekonomi	0	1	40	38	1
pusing	4	5	1	3	30
keuangan	1	2	30	25	2
sakit	4	6	0	4	25
Inflasi	8	1	15	14	1

vector of d_3
 $= [40, 1, 30, 0, 15]$

- Two documents are similar if they have similar vector!

$$d_3 = [40, 1, 30, 0, 15]$$

(看的是列向量)

$$d_4 = [38, 3, 25, 4, 14]$$

Vector Space Model



- **Term-Document Matrix**

Each cell is the count of word t in document d

	d_1	d_2	d_3	d_4	d_5
ekonomi	0	1	40	38	1
pusing	4	5	1	3	30
keuangan	1	2	30	25	2
sakit	4	6	0	4	25
Inflasi	8	1	15	14	1

Vector of the word
“sakit” = $[4, 6, 0, 4, 25]$

- Two words are similar if they have similar vector!

pusing = $[4, 5, 1, 3, 30]$

(看的是行向量)

sakit = $[4, 6, 0, 4, 25]$

Vector Space Model



- **Weighting:** in practice, we usually use weights such as **TF-IDF**, instead of just using **raw counts (only TF)**.

$$\text{tf-idf}_{w,d} = \text{tf}_{w,d} \times \log(N / \text{df}_w)$$

$\text{tf}_{w,d}$ = frequency of w in d
 df_w = number of documents containing w
 N = total number of documents

log项是一个惩罚项, 包含单词 w 的文档数量越多, 同一个文档对于单词 w 的贡献越小, 这样的话更容易体现出单词的独特性, 即只出现在少数文档中的单词应该有更大的count。

	d_1	d_2	d_3	d_4	d_5
ekonomi	0	1	40	38	1
pusing	4	5	1	3	30
keuangan	1	2	30	25	2
sakit	4	6	0	4	25
Inflasi	8	1	15	14	1

$$\text{TF}(\text{sakit}) = [4, 6, 0, 4, 25]$$

$$\left. \begin{array}{l} \text{DF}(\text{sakit}) = 4 \\ N = 5 \end{array} \right\} \text{IDF}(\text{sakit}) = \log(5/4)$$

$$\text{TF-IDF}(\text{sakit}) = [\dots\dots]$$

Limitations of Vector Space Model



- TF-IDF vectors are
 - **long** (length $|V| = 20,000$ to $50,000$)
 - **sparse** (**most elements are zero**)
- difficult to use as features in machine learning (more weights to tune)
- storing explicit counts can be difficult for generalization



Prediction based: word2vec

Word2Vec: Overview



- **Word2vec** (Mikolov et al. 2013) is a framework for learning word vectors.

Idea:



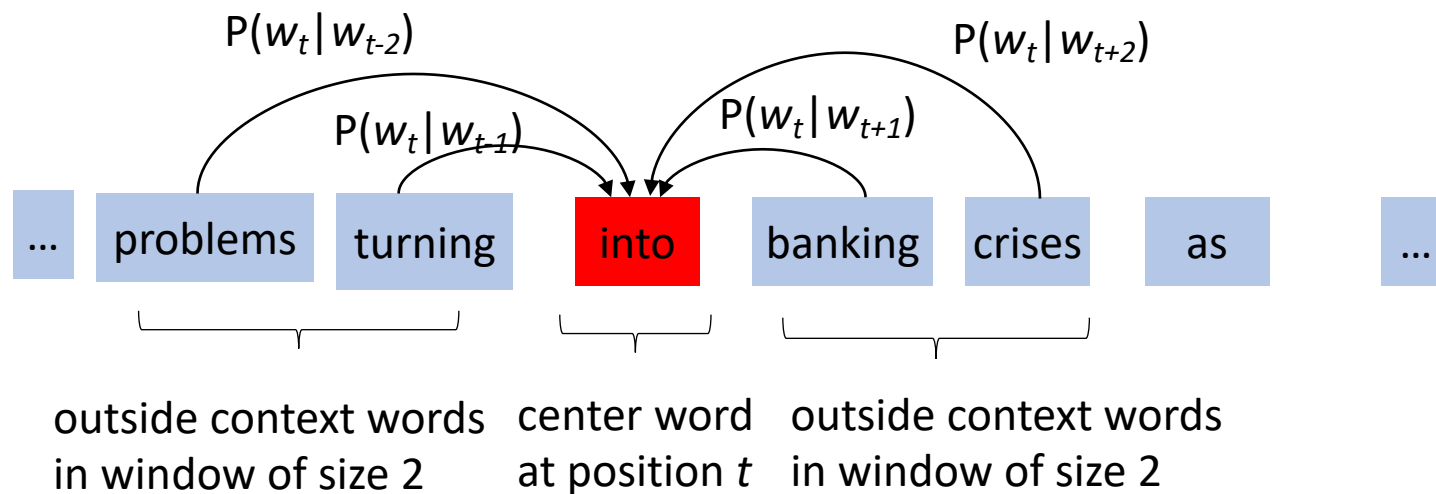
- we have a **large corpus of text**.
- every word in a fixed vocabulary is represented by a **learnable vector**;
- go through each position t in the text, which has a **center word** c and **context** (“outside”) words o ;
- use the **similarity of the word vectors** for c and o to **calculate the probability** of o given c (or vice versa) ;
- keep **adjusting** the word vectors to maximize this probability.

Mikolov T, Sutskever I, Chen K, et al. Distributed representations of words and phrases and their compositionality. NIPS 2013.

Word2Vec: Overview



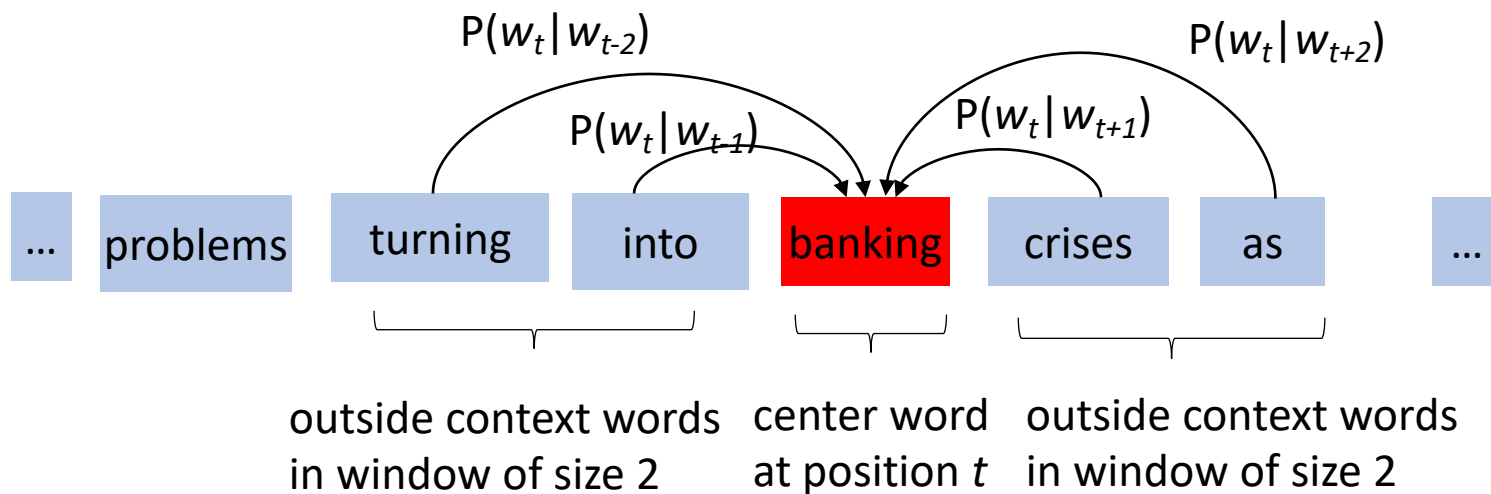
Example: **window** and process for computing $P(w_t | w_{t+j})$



Word2Vec: Overview



Example: window and process for computing $P(w_t | w_{t+j})$



Mikolov's CBOW

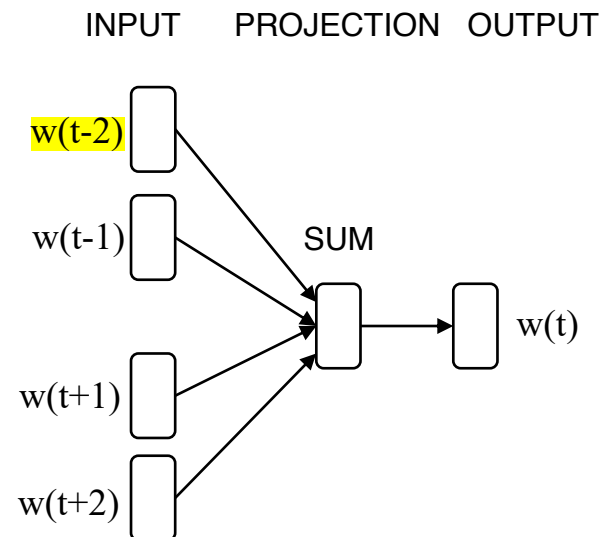


- **CBOW**: the distributed representations of context (or surrounding words) are combined to **predict the word in the middle**.

$$P(w_t | w_{t-k}, \dots, w_{t-1}, w_{t+1}, \dots, w_{t+k}) = \text{softmax}(\mathbf{NN}(V(w_{t-k}) + \dots + V(w_{t+k})))$$

This can be represented by a **neural network**:

- An **input layer** which converts each **word (one-hot)** into a **dense vector**.
- A **projection layer** which **combines the vectors** of input words.
- An **output layer** which predicts the target word w_t given the combined context vector.



CBOW = **Continuous**-Bag-of-Words continuous指的是通过上下文中连续的几个单词来判断中心单词的语义。

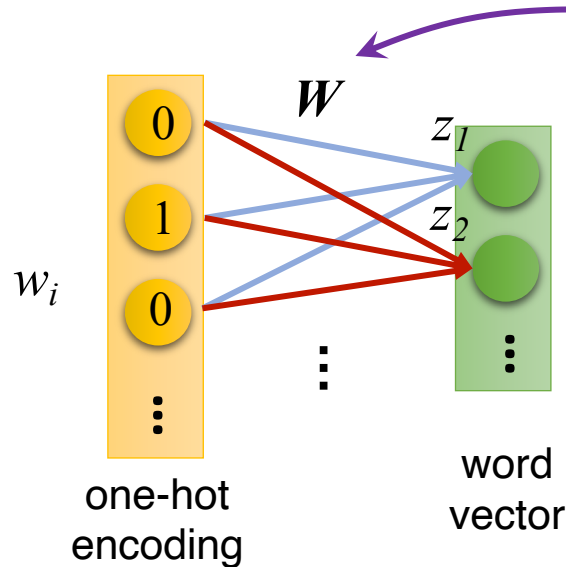
the order of words in the context does not influence the projection.

Model Architecture

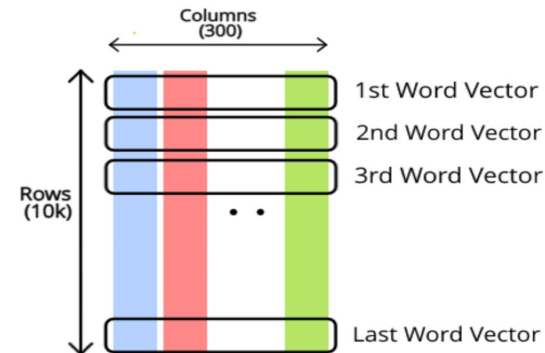


- Input Layer: represent any word into a vector.

$$z_i = Wx_i = W_i$$



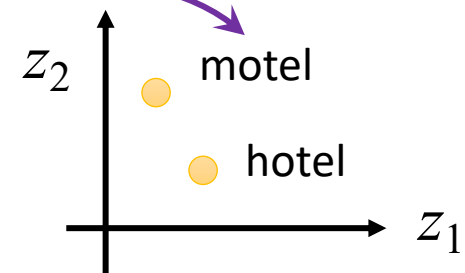
The weight matrix $W \in R^{|V| \times d}$ is a **lookup table** with each row W_i being the vector for word w_i .



Example

$$[0 \ 0 \ 0 \ 1 \ 0] \times \begin{bmatrix} 17 & 24 & 1 \\ 23 & 5 & 7 \\ 4 & 6 & 13 \\ 10 & 12 & 19 \\ 11 & 18 & 25 \end{bmatrix} = [10 \ 12 \ 19]$$

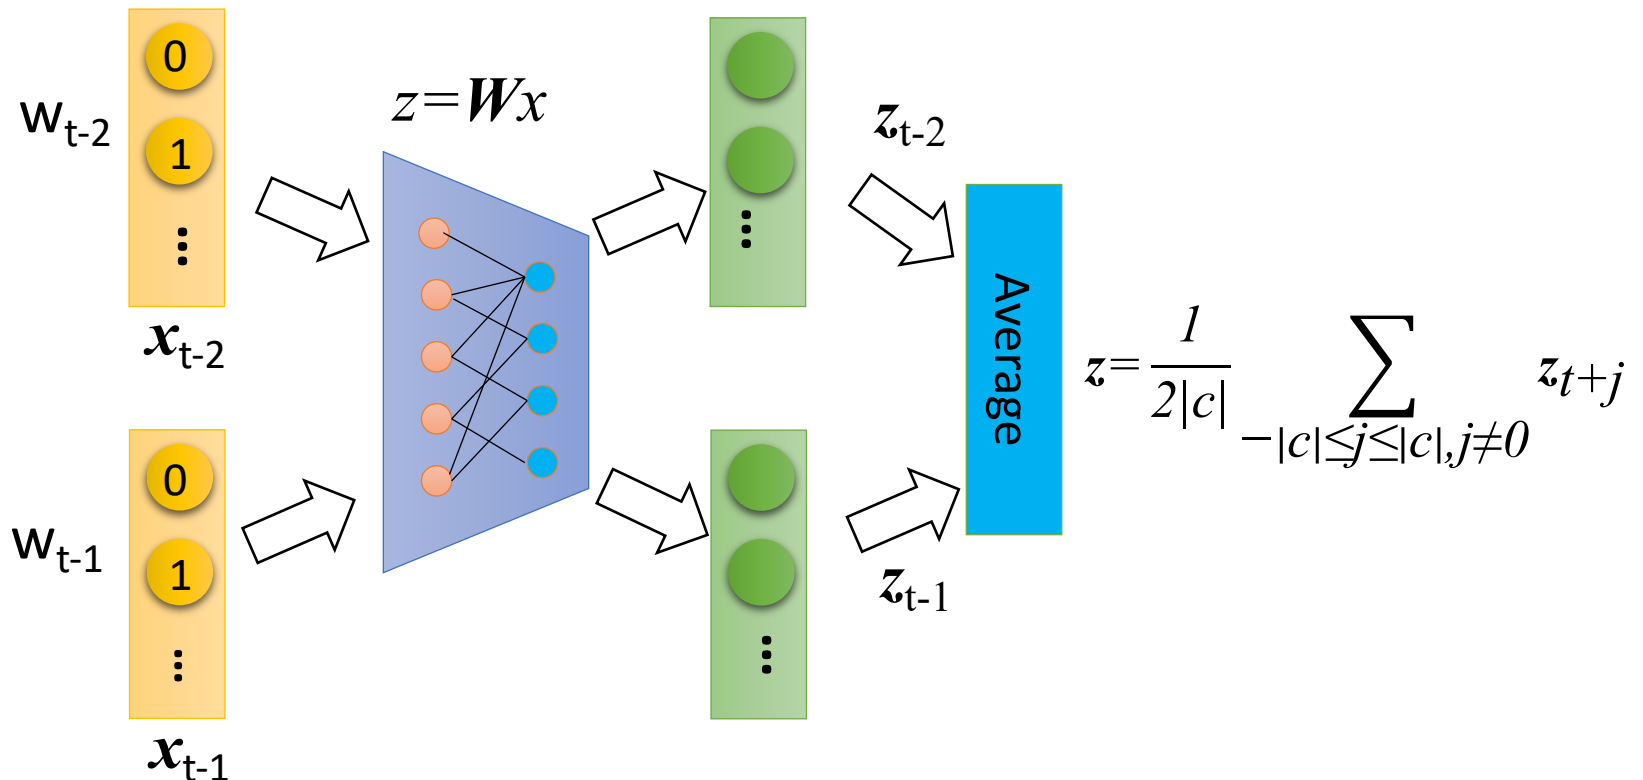
"motel"



Architecture



- Projection Layer: combining context vectors into one vector.

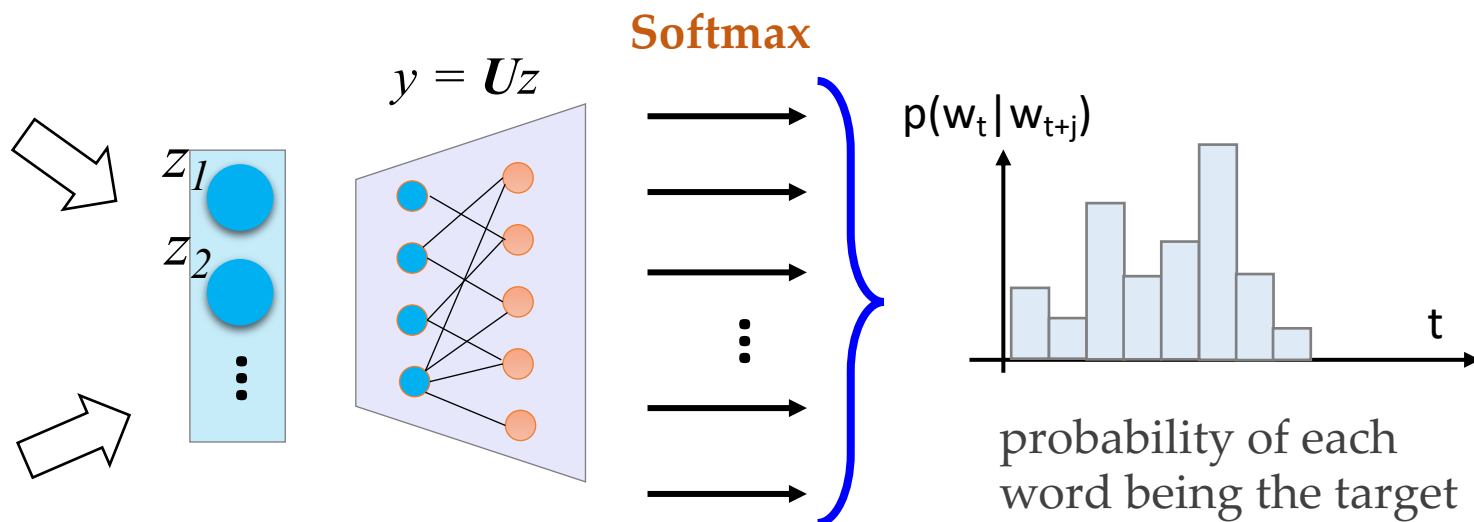


Architecture

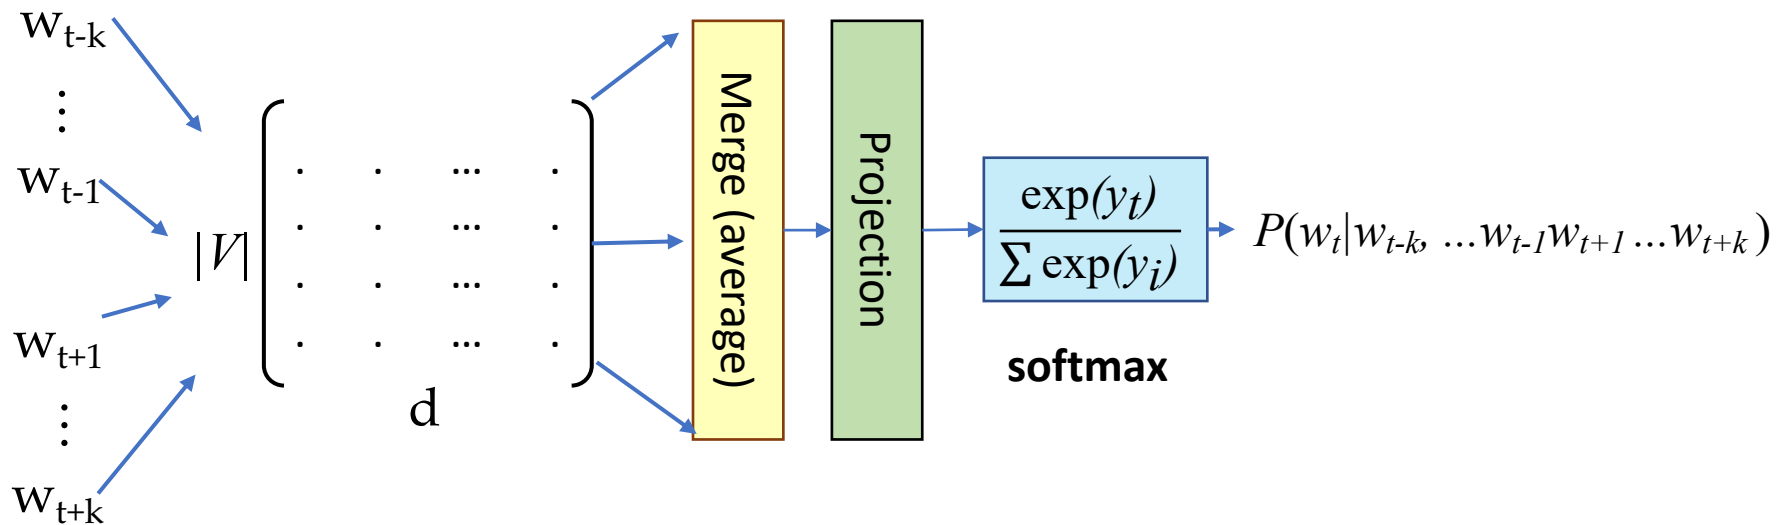


- Output Layer: predicts the probability of the target word.

$$P(w_t | w_{t-|c|}, \dots, w_{t-1}, w_{t+1}, \dots, w_{t+|c|}) = \frac{\exp(y_t)}{\sum_{i=1}^{|V|} \exp(y_i)}$$



Architecture: the big picture



Training



- Given $D = \{w_1, w_2, \dots, w_N\}$, minimize the **negative log likelihood (NLL)** loss function:

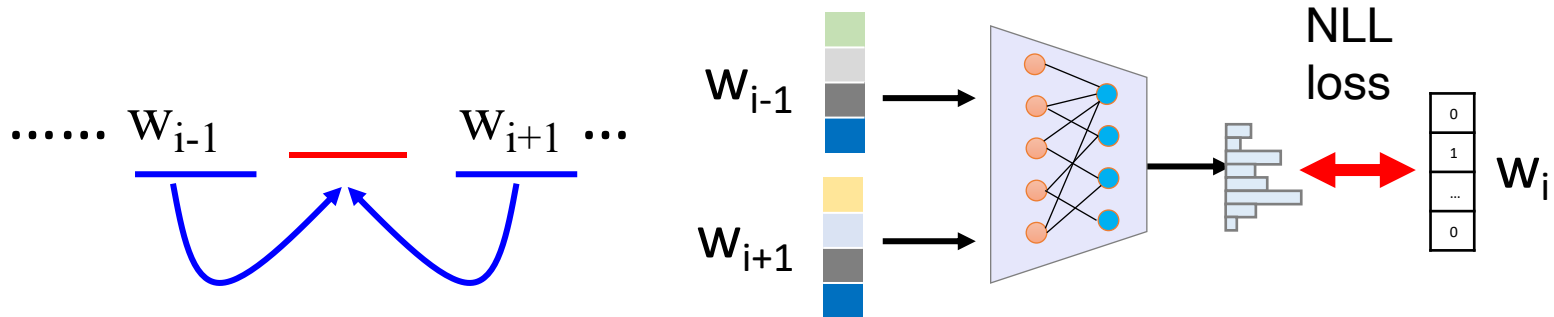
$$L(W, U | D) = -\frac{1}{N} \sum_{t=1}^N \log p(w_t | w_{t-k}, \dots, w_{t-1}, w_{t+1}, \dots, w_{t+k})$$

using **gradient descend**.

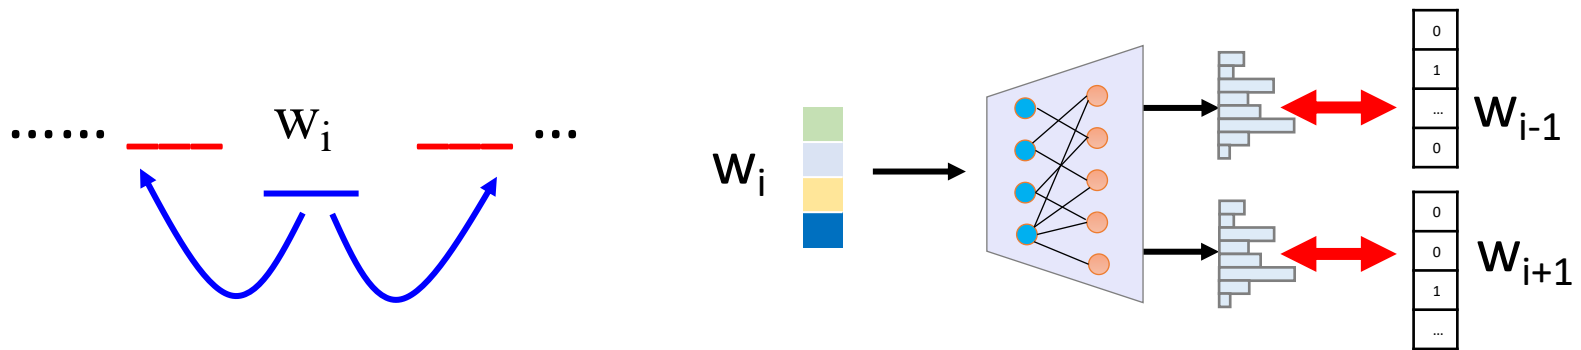
Other Models



- CBOW predict the word given its context



- Skip-gram predict the context given a word



The Skip-Gram Model

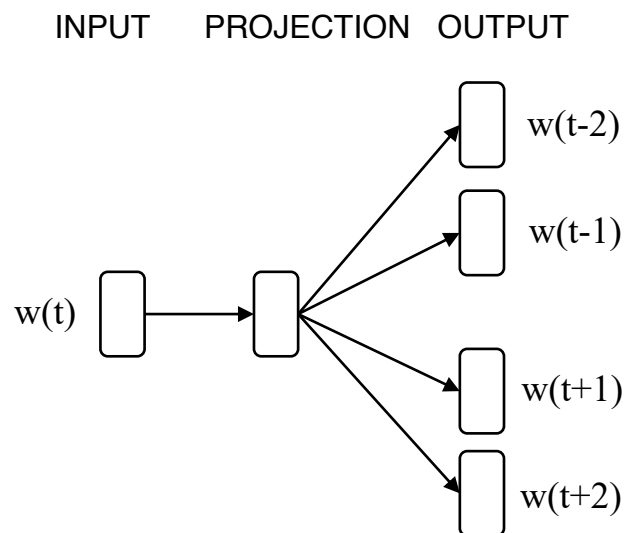


- We seek a model for $P(w_{t+j} | w_t)$.

$$P(w_{t+j} | w_t) = \frac{\exp(y_{t+j})}{\sum_{i=1}^{|V|} \exp(y_i)}$$

$$y = U z$$

$$z = W x$$

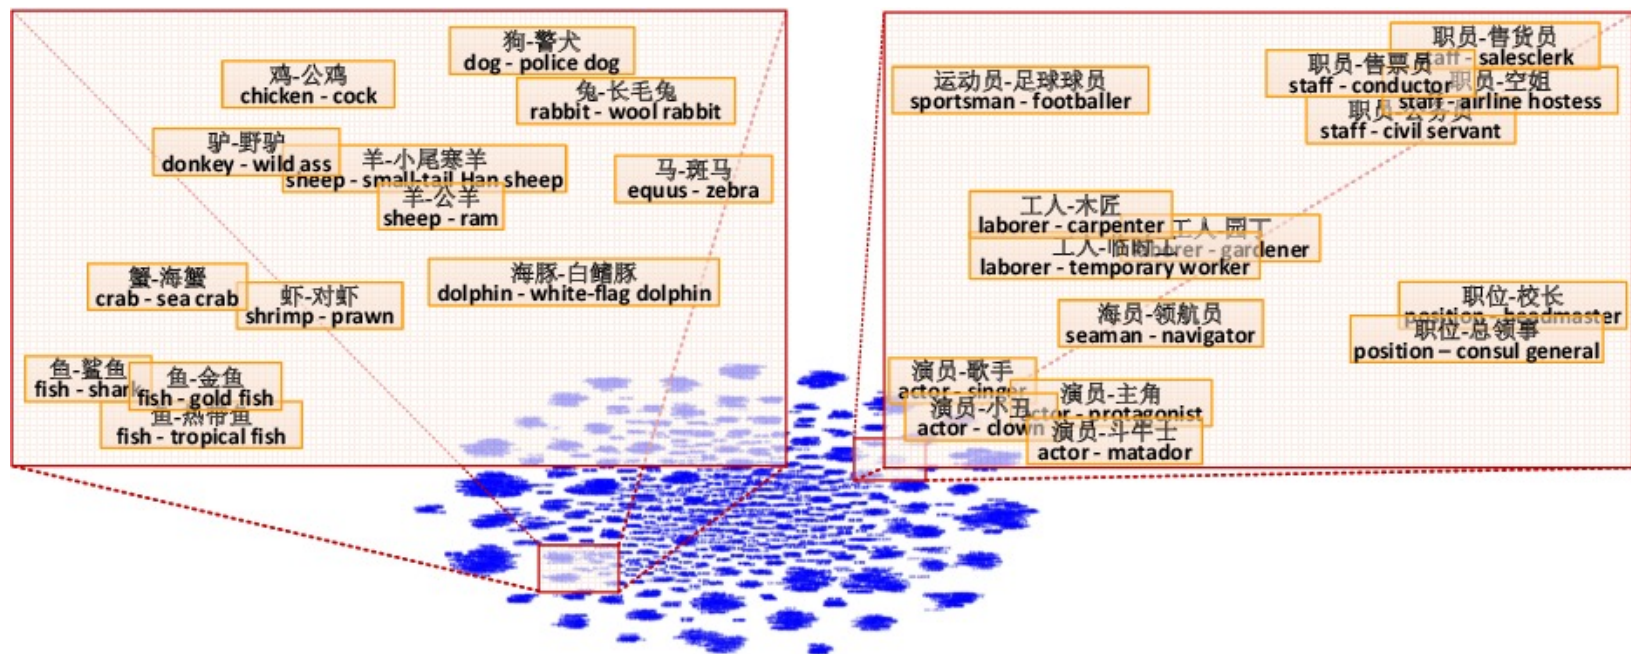


$$L(W, U | \chi) = -\frac{1}{N} \sum_{t=1}^N \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

Example



- Word2vec maximizes an objective function by putting similar words nearby in space.



Fu, Ruiji, et al. "Learning semantic hierarchies via word embeddings." *ACL* 2014.

The Word Analogy Task analogy : 类比



- Word Analogy:

a:b :: c:?

man:woman :: king:?

Examples

- Man is to Woman as King is to ___?
- Good is to Best as Smart is to ___?
- China is to Beijing as America is to ___?

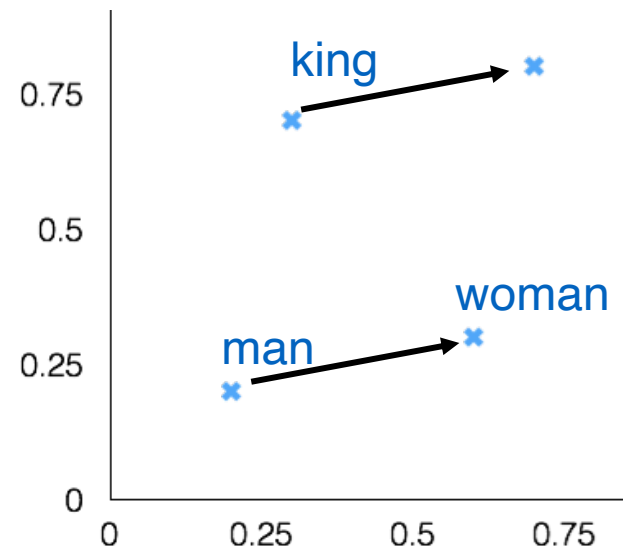
How to find d ?

$$d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{||x_b - x_a + x_c||}$$

- It turns out that word2vec is good for such an analogy task.

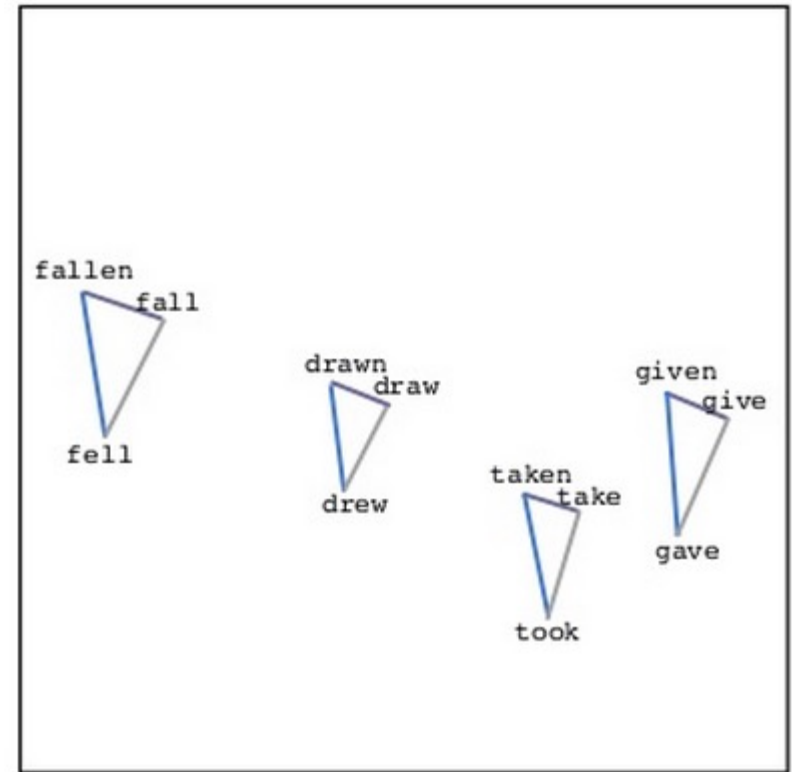
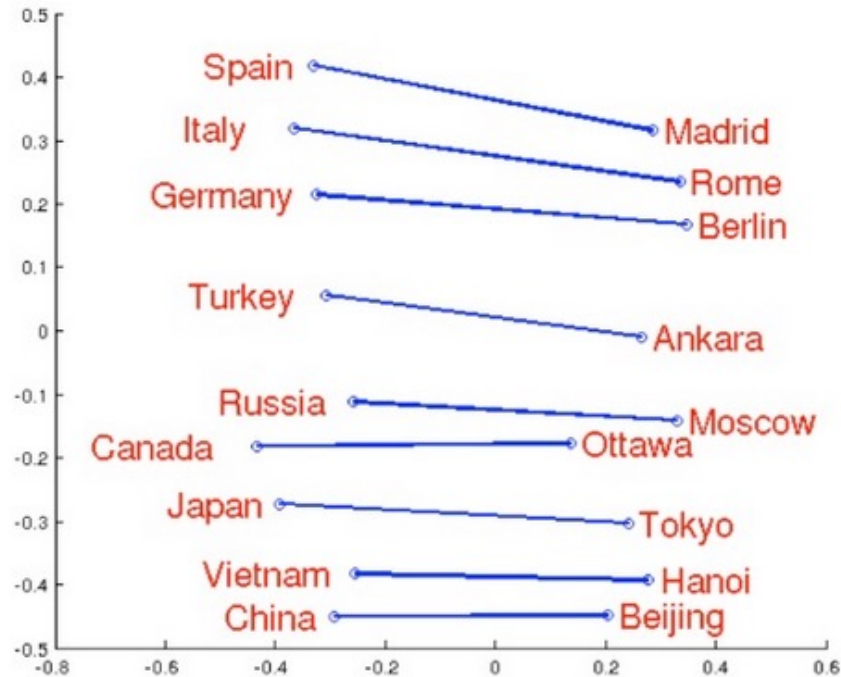
$$V_{\text{king}} - V_{\text{man}} + V_{\text{woman}} = V_{\text{queen}}$$

平行且同向时显然最大



Mikolov, T., et al. Distributed Representations of Words and Phrases and Their Compositionality. NIPS 2013.

Example



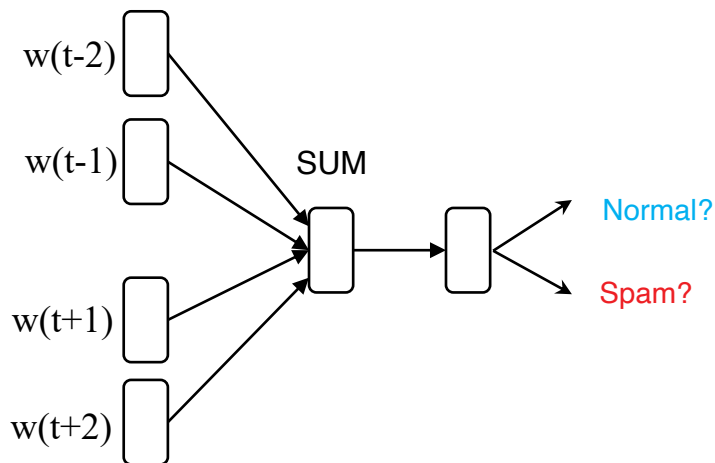
Source: <http://www.slideshare.net/hustwj/cikm-keynotenov2014>

Thinking



总体而言是一个非监督学习,但是其中间过程是需要softmax的,是一个分类模型,所以其中间过程也包括一部分的监督学习

- What kind of machine learning is word2vec?
- What limitations does word2vec have if we use it for text classification?

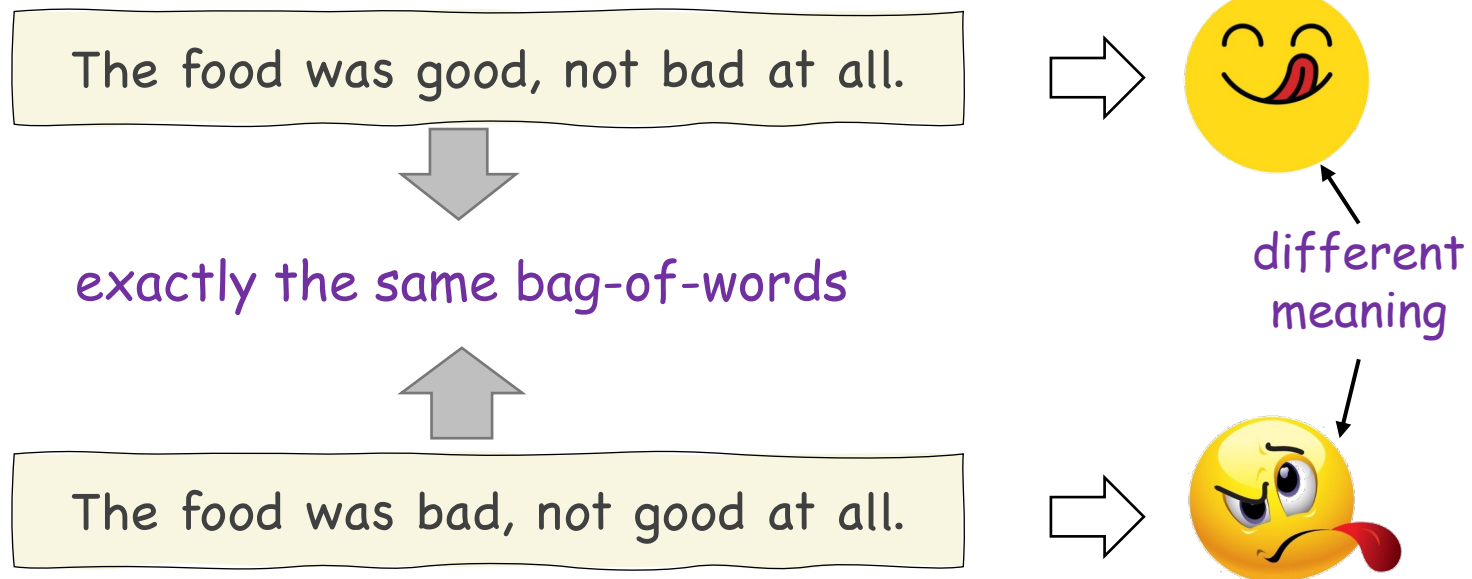


Beyond Bag-of-Words

Order
Matters!



- To understand the meaning of a sentence, the **order** of the words can not be ignored.



Paragraph Vector: Le, Quoc, and Tomas Mikolov. "Distributed Representations of Sentences and Documents." ICML, 2014

Seq2seq Auto-encoder: Li, Jiwei, Minh-Thang Luong, and Dan Jurafsky. "A hierarchical neural autoencoder for paragraphs and documents." arXiv preprint, 2015

Skip Thought: Ryan Kiros, Yukun Zhu, Ruslan Salakhutdinov, Richard S. Zemel, Antonio Torralba, Raquel Urtasun, Sanja Fidler, "Skip-Thought Vectors" arXiv preprint, 2015.

Beyond Bag-of-Words



- **Language Models**

Recurrent Neural Networks, Transformers,...

- **Pretrained Language Models**

BERT, GPT-2, ...

To appear in future lectures..

Language Models



- A probabilistic model of **how likely** a **given string appear in a given “language”**.
- For any sequence $x = (w_1, w_2, \dots, w_N)$, a **language model** can be defined as:

$$p(x) = p(w_1, w_2, \dots, w_N)$$

由于这个语言模型是判断一整个string出现在language的概率，所以其研究的应该是各个string中单次同时出现的概率，所以使用的是联合概率分布。

Example:

$P_1 = P(\text{“我爱机器学习”})$
 $P_2 = P(\text{“我爱学习机器”})$
 $P_3 = P(\text{“机器我爱学习”})$
 $P_4 = P(\text{“爱我机学习器”})$

Chinese: $P_1 > P_2 > P_3 > P_4$

• Applications:

message suggestion; document generation; spelling correction; machine translation; speech recognition;...

我爱机器学__?

Language Model



- What is the probability of $P(w_1, \dots, w_N)$?

$$p(\text{我爱机器学习}) = ?$$

- Chain Rule:

$$p(w_1, \dots, w_N) = p(w_1)p(w_2|w_1)\dots p(w_N|w_1, \dots, w_{N-1})$$

$$p(\text{我爱机器学习}) = p(\text{我})p(\text{爱}|\text{我})p(\text{机}|\text{我爱})p(\text{器}|\text{我爱机})p(\text{学}|\text{我爱机器})p(\text{习}|\text{我爱机器学})$$

- Markov Assumption: (only consider the last $n-1$ words)

$$p(w_i|w_1, \dots, w_{i-1}) = p(w_i|w_{i-n+1}, \dots, w_{i-1})$$

$$p(\text{习}|\text{我爱机器学}) \approx p(\text{习}|\text{机器学}) \approx p(\text{习}|\text{学})$$

Bigram Language Model



So that's what we get for **n=2**:

$$p(w) = p(w_1)p(w_2|w_1)\dots p(w_N|w_{N-1})$$

$$1/18 \times 1/8 \times 1/120 \times 1/4 \times 1/420 \times 1/2$$

$$p(\text{我爱机器学习}) = p(\text{我})p(\text{爱}|\text{我})p(\text{机}|\text{爱})p(\text{器}|\text{机})p(\text{学}|\text{器})p(\text{习}|\text{学})$$

Q: How to estimate these probabilities?



A: Straightforward counting.

Q: But remember in word embedding, counting discrete words have many drawbacks?

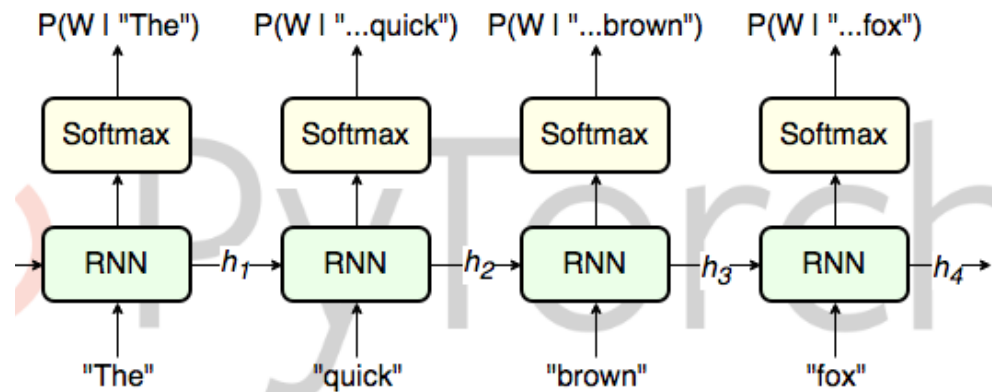
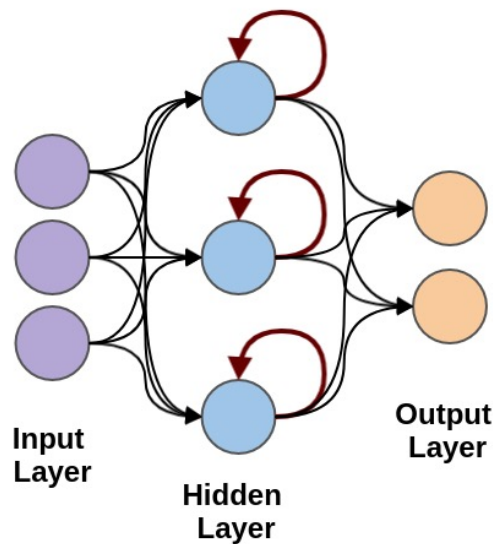
A: Don't worry, we have **neural networks as language models**.

What's Next?

WHAT'S
NEXT?

Recurrent Neural Networks

- A deep neural network for sequence (language) modeling.



NEXT